



Faculty of Science



The Learning Problem

Statistical Methods for Machine Learning

Christian Igel
Department of Computer Science



Machine learning

Machine learning is a branch of computer science and applied statistics covering software that improves its performance at a given task based on sample data or experience.



Why machine learning?

- Computer systems are required for tasks for which solutions cannot be specified in the traditional way, e.g., because
 - the designer's knowledge is limited, and/or
 - the sheer complexity and variability precludes an accurate description.
- However, large amounts of data describing the task are often available or can be automatically obtained.
- To take proper advantage of this information, we need systems that self-adapt and automatically improve based on sample data – systems that learn.
- Machine learning (ML) is a vital ingredient of data mining, data analytics, knowledge discovery, etc.

Machine learning turns data into knowledge



Outline

- ① Supervised Learning Scenario
- ② Nearest Neighbor Classification
- ③ Cross-validation for Hyperparameter Selection
- ④ Generalization and Regularization
- ⑤ Summary



Outline

- ① Supervised Learning Scenario
- ② Nearest Neighbor Classification
- ③ Cross-validation for Hyperparameter Selection
- ④ Generalization and Regularization
- ⑤ Summary



Inductive inference

Learning is goal-directed changing of behaviour based on experience.

Supervised machine learning formalizes and automates the process of *inductive inference*:

- 1 Observe a phenomenon
- 2 Construct a model of that phenomenon
- 3 Make predictions using this model



Binary classification

two classes: {apples, pears}, $\{-1, 1\}$




Supervised learning

The goal of supervised learning is to uncover an unknown relation between

- input space \mathcal{X} ,
e.g., real-valued vector describing color and shape of fruit, and
- output space \mathcal{Y} ,
e.g., {apples, pears},

based on sample input-output data $S = \{(x_1, y_1), \dots, (x_\ell, y_\ell)\}$.




I.i.d. assumption

- *I.i.d.* (identically, independently distributed) assumption: Each pattern is drawn independently from an (unknown) stable distribution p over $(\mathcal{X}, \mathcal{Y})$:

$$p((x, y)) = p(x) \cdot p(y | x)$$



- Example: Probability of (red, 8 cm, round) is always the same for each training or test pattern

Probability of being an apple is the same for each fruit that is (red, 8 cm, round)



The goal of learning

Given sample data, we want to find a hypothesis predicting output from input showing good average performance on patterns $(x_i, y_i) \sim p$.



Broad division of supervised learning

- ① Classification: Assigning an input to one class of a finite set of classes $\mathcal{Y} = \{\mathcal{C}_1, \dots, \mathcal{C}_m\}$, $2 \leq m < \infty$.
- ② Regression: Given an input, predict corresponding output from $\mathcal{Y} \subseteq \mathbb{R}^m$, $1 \leq m < \infty$.
- ③ Density estimation: Given an input x , predict the probability distribution $p(y|x)$ on \mathcal{Y} .

We focus on classification and regression in the following.



Learning machines

- A *hypothesis* $h : \mathcal{X} \rightarrow \mathcal{Y}$ maps inputs to outputs, a hypothesis class \mathcal{H} is a set of such functions
E.g., a hypothesis makes a prediction whether a fruit is an apple or a pear given shape, color, and size
- Applying a learning algorithm means coming up with a hypothesis given sample data (formally, it maps from $\{((x_1, y_1), \dots, (x_\ell, y_\ell)) \mid 1 \leq i \leq \ell < \infty, x_i \in \mathcal{X}, y_i \in \mathcal{Y}\}$ to \mathcal{H}).



Loss/risk functions

The quality of the prediction of a hypothesis is quantified by a *task-dependent* loss function.

A function $L : \mathcal{Y} \times \mathcal{Y} \rightarrow [0, \infty[$ with the property $L(y, y) = 0$ for $y \in \mathcal{Y}$ is called a *loss function*.

Typical loss for classification is the 0-1 loss:

$$L(y, \hat{y}) = \begin{cases} 0 & \text{if } y = \hat{y} \\ 1 & \text{otherwise} \end{cases}$$



Goal of learning

Goal: Given sample data, we want to find a hypothesis h showing good average performance on patterns $(x, y) \sim p$.

That is, we want to find h minimizing the *risk*:

$$\mathcal{R}_p(h) = \mathbb{E}_{(x,y) \sim p} \{L(y, h(x))\}$$

The *empirical risk* of h on sample data S ($|S| = \ell$) is defined as:

$$\mathcal{R}_S(h) = \frac{1}{\ell} \sum_{i=1}^{\ell} L(y_i, h(x_i))$$



Typical loss functions

- **Classification:** Under the 0-1 loss $L(y, h(x)) = \mathbb{I}\{h(x) \neq y\}$ the risk of hypothesis h is the probability of error

$$\mathcal{R}_p(h) = \Pr(h(X) \neq Y) = \mathbb{E}_p\{\mathbb{I}\{h(X) \neq Y\}\}$$

and the empirical risk on sample data S ($|S| = \ell$) counts errors:

$$\mathcal{R}_S(h) = \frac{1}{\ell} \sum_{i=1}^{\ell} \mathbb{I}\{h(x_i) \neq y_i\}$$

(indicator function $\mathbb{I}\{\cdot\}$ is 1 if its argument is true and 0 otherwise)

- **Regression:** Under the squared error

$$L(y, \hat{y}) = (y - \hat{y})^2$$

the empirical risk corresponds to the sum-of-squares error



Bayes risk and consistency

- Minimum risk over *all possible measurable functions* h is the Bayes risk

$$\mathcal{R}_p^{\text{Bayes}} = \mathcal{R}_p(h^{\text{Bayes}}) = \inf_h \mathcal{R}_p(h)$$

usually differs from $\mathcal{R}_p^* = \mathcal{R}_p(h^*) = \inf_{h \in \mathcal{H}} \mathcal{R}_p(h)$

for $\mathcal{Y} = \{-1, 1\}$ we have

$$h^{\text{Bayes}}(x) = \text{sgn} [\mathbb{E}_p(Y \mid X = x)]$$

(here $\text{sgn}(x) := 2\mathbb{I}\{x > 0\} - 1$)

- Algorithm a is *consistent* if $\lim_{\ell \rightarrow \infty} \mathcal{R}_S(a(S)) = \mathcal{R}_p^{\text{Bayes}}$ almost surely




Outline

- ① Supervised Learning Scenario
- ② Nearest Neighbor Classification
- ③ Cross-validation for Hyperparameter Selection
- ④ Generalization and Regularization
- ⑤ Summary





Distinguishing stars and galaxies



A typical classification task

L. K. Polsterer, P.-C. Zinn, F. Gieseke. Finding New High-Redshift Quasars by Asking the Neighbours. *Monthly Notices of the Royal Astronomical Society (MNRAS)* 428(1):226–235, 2013



Nearest neighbor classification

- Idea:
 - Given training data $S = \{(x_1, y_1), \dots, (x_\ell, y_\ell)\}$ and new data point x
 - Look in S for x_i most similar to x
 - Assign new record to class of x_i
- This requires a notion of similarity



Recall: Metric

Definition: A *metric* or *distance function* on \mathcal{X} is a function $d : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ with:

- ① $d(x, y) \geq 0$ with $d(x, y) = 0$ if and only if $x = y$,
 - ② $d(x, y) = d(y, x)$ (symmetry),
 - ③ $d(x, z) \leq d(x, y) + d(y, z)$ (triangle inequality)
- for all $x, y \in \mathcal{X}$.

Example: Euclidean metric

$$d(\mathbf{p}, \mathbf{q}) = \sqrt{(p_1 - q_1)^2 + (p_2 - q_2)^2 + \cdots + (p_n - q_n)^2}$$

for $\mathbf{p} = (p_1, \dots, p_n)^\top$, $\mathbf{q} = (q_1, \dots, q_n)^\top \in \mathbb{R}^n$



1-Nearest neighbor (1-NN) classification

Algorithm 1: 1-nearest neighbor


Input: metric $d : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$, training data $S = \{(x_1, y_1), \dots\} \in (\mathcal{X} \times \{-1, 1\})^\ell$, new input x to be classified

Output: predicted label y of x

/* ties are broken at random */

1 $(x_{\min}, y_{\min}) = \operatorname{argmin}_{(x_i, y_i) \in S} d(x_i, x)$

Result: $y = y_{\min}$



K -Nearest neighbor (K -NN) classification

Algorithm 2: K -nearest neighbor

Input: metric d , training data

$$S = \{(x_1, y_1), \dots, (x_\ell, y_\ell)\},$$


input x to be classified, number
of neighbors K

Output: predicted label y of x


```
1 while  $|S^*| < K$  do
  /* ties broken at random */
2    $(x_{\min}, y_{\min}) = \operatorname{argmin}_{(x_i, y_i) \in S} d(x_i, x)$ 
3    $S \leftarrow S \setminus \{(x_{\min}, y_{\min})\}$ 
4    $S^* \leftarrow S^* \cup \{(x_{\min}, y_{\min})\}$ 
```

```
/* ties broken at random */
```

Result: $y = \operatorname{argmax}_c |\{(x, y) \mid (x, y) \in S^* \wedge y = c\}|$



Influence of K



C. M. Bishop. *Pattern Recognition and Machine Learning*, Springer, 2006



Outline

- ① Supervised Learning Scenario
- ② Nearest Neighbor Classification
- ③ Cross-validation for Hyperparameter Selection
- ④ Generalization and Regularization
- ⑤ Summary




Hyperparameter selection

- The results of K -NN depends on the number of neighbors – now, how choose K ?
- Algorithm parameters such as K are called *hyperparameters*.
- Finding hyperparameters is often called *model selection*.
- A standard method for selection of a few hyperparameters is κ -fold *cross-validation*.



Cross-validation



C. M. Bishop. *Pattern Recognition and Machine Learning*, Springer, 2006

- Split S into κ (almost) even parts
- Train κ times on $\kappa - 1$ parts and test on the held-out part
- Take hyperparameters maximizing mean test performance for building the model using all data





Outline

- 1 Supervised Learning Scenario
- 2 Nearest Neighbor Classification
- 3 Cross-validation for Hyperparameter Selection
- 4 Generalization and Regularization**
- 5 Summary



Learning by heart is not enough

- 1-NN has always zero error on the training set. And on a test set?
- Minimizing empirical risk does not imply good *generalization* (i.e., good performance on patterns not seen during training).
- *Overfitting*: The hypothesis faithfully reflects idiosyncrasies of the training data rather than the underlying distribution.



Learning strategies

- Empirical risk minimization given \mathcal{H}
("learning by heart restricted to \mathcal{H} ")

$$h_S = \operatorname{argmin}_{h \in \mathcal{H}} \mathcal{R}_S(h)$$

- Regularized risk minimization

$$h_S = \operatorname{argmin}_{h \in \mathcal{H}} [\mathcal{R}_S(h) + \text{penalty}(h)]$$



Examples of regularization

- Using the norm for regularization

$$h_S = \operatorname{argmin}_{h \in \mathcal{H}} [\mathcal{R}_S(h) + \gamma \|h\|] \quad (\gamma > 0)$$

- Bayesian normalization

$$p(h | S) = \frac{p(S | h)p(h)}{p(S)} \propto p(S | h)p(h)$$

$$\text{posterior} \propto \text{likelihood} \times \text{prior}$$

instead of maximizing the a posteriori probability (MAP) it is convenient to minimize its negative logarithm

$\text{likelihood} \propto \exp(-\gamma \mathcal{R}_S(h))$ yields for MAP

$$h_S = \operatorname{argmin}_{h \in \mathcal{H}} [\gamma \mathcal{R}_S(h) - \log \text{prior}(h)]$$



Early stopping

- There are ways of avoiding overfitting not falling in the general regularization framework outlined before.
- *Early-stopping*: The learning algorithm
 - partitions sample S into training S_{train} and validation S_{val} data,
 - produces iteratively a sequence of hypotheses


$$h_1, h_2, h_3, \dots$$

based on S_{train} , ideally corresponding to a nested sequence of hypothesis spaces $\mathcal{H}_1 \subseteq \mathcal{H}_2 \dots$ with $h_i \in \mathcal{H}_i$ and

- non-decreasing complexity and
- decreasing empirical risk $\mathcal{R}_{S_{\text{train}}}(h_i) > \mathcal{R}_{S_{\text{train}}}(h_{i+1})$ on S_{train} ,
- monitors empirical risk $\mathcal{R}_{S_{\text{val}}}(h_i)$ on the validation data, and
- outputs the hypothesis h_i minimizing $\mathcal{R}_{S_{\text{val}}}(h_i)$.



Early stopping example



Errors & bounds

Typical bound:


With probability of at least $1 - \delta$

$$\mathcal{R}_p(h_S) \leq \mathcal{R}_S(h_S) + B(\ell, \delta, \mathcal{H})$$

Carefully check assumptions and to which entities expectations refer!



Free lunch...



No free lunch

- ① If there is no assumption how the past (training data) is related to the future (test data), prediction is impossible.
- ② If there is no a priori restriction on the possible phenomena that are expected, it is impossible to generalize and thus no algorithm is superior to another.
- ③ Any consistent algorithm can have arbitrarily bad behavior when given a finite, incomplete training set.

Generalization = Data + Knowledge


OLIVIER BOUSQUET

O. Bousquet, S. Boucheron, G. Lugosi. Introduction to Statistical Learning Theory. In *Advanced Lectures in Machine Learning*, LNAI 3176, 2004



Prior knowledge

But in real-world applications we can make assumptions about the expected phenomena!



Prior knowledge can be incorporated by

- a proper restriction of the hypothesis space,
- a proper definition of neighbourhood/similarity (e.g., by defining a proper scalar product inducing a metric), and/or
- penalising the “complexity” of classifiers (i.e., preferring “simple” solutions).



Outline

- ① Supervised Learning Scenario
- ② Nearest Neighbor Classification
- ③ Cross-validation for Hyperparameter Selection
- ④ Generalization and Regularization
- ⑤ Summary



Summary

- The goal of supervised learning is to uncover an unknown relation between an input and an output space given examples.
- Memorizing training examples is not enough to achieve generalization (i.e., to classify correctly unseen examples).
- Incorporation of prior knowledge is required for generalization, for example by restricting the hypotheses space, by carefully designed distance measures, and/or by penalties on solution complexity.
- Nearest neighbor classification is a simple algorithm giving good (baseline) results in practice.
- Cross-validation can be used to find hyperparameters.

